

Binary Classification of Pulsar Detection

Francesco Sorrentino
Polytechnic of Turin
Artificial Intelligence and Data Analytics
s301665

Francesco Di Gangi
Polytechnic of Turin
Artificial Intelligence and Data Analytics
s301793

Abstract

In this report is shown how some of the most common Machine Learning (ML) algorithms in order to classify and compare the results of the HTRU2 dataset. The starting point will be the feature analysis and feature correlation, then it will show the result of different classifiers using normalization techniques (Gaussian and Z Score normalization) in order to lead us to choose the best-fit model.

1 Problem Overview

The data set used in this project is called *HTRU2*, and it describes a sample of pulsar candidates collected during the High Time Resolution Universe Survey (South) [1]. Pulsars are a rare type of Neutron star that produce radio emission detectable here on Earth. They are of considerable scientific interest as probes of space-time, the inter-stellar medium, and states of matter.

In this report it is possible to see how **Machine Learning** tools are now being used automatically label pulsar candidates to facilitate rapid analysis, especially in the *classification* field.

This report is about binary classification problems, reason why the data set is divided in two classes:

- 1,639 legitimate pulsar examples are a minority positive class, real pulsar examples
- 16,259 spurious examples the majority negative class, caused by RFI/noise

for a total of 17,898 examples, that are split into training set (8929 samples) and test set (8969 samples).

Also, each candidate is described by eight continuous variables:

1. Mean of the integrated profile.
2. Standard deviation of the integrated profile.
3. Excess kurtosis of the integrated profile.
4. Skewness of the integrated profile.
5. Mean of the DM-SNR curve.
6. Standard deviation of the DM-SNR curve.
7. Excess kurtosis of the DM-SNR curve.
8. Skewness of the DM-SNR curve.

The first four are simple statistics obtained from the integrated pulse profile (folded profile), that is an array of continuous variables describing a longitude-resolved version of the signal that has been averaged in both time and frequency, meanwhile the remaining four variables are similarly obtained from the DM-SNR curve.

The data are normalized following the *Z-Score Normalization*

$$z = \frac{x_i - \mu}{\sigma}, \forall i \in \{1, \dots, n\}$$

where x_i indicates the observed value for the i -th sample, and μ is the mean vector of samples, ending with *sigma* that represents the standard deviation.

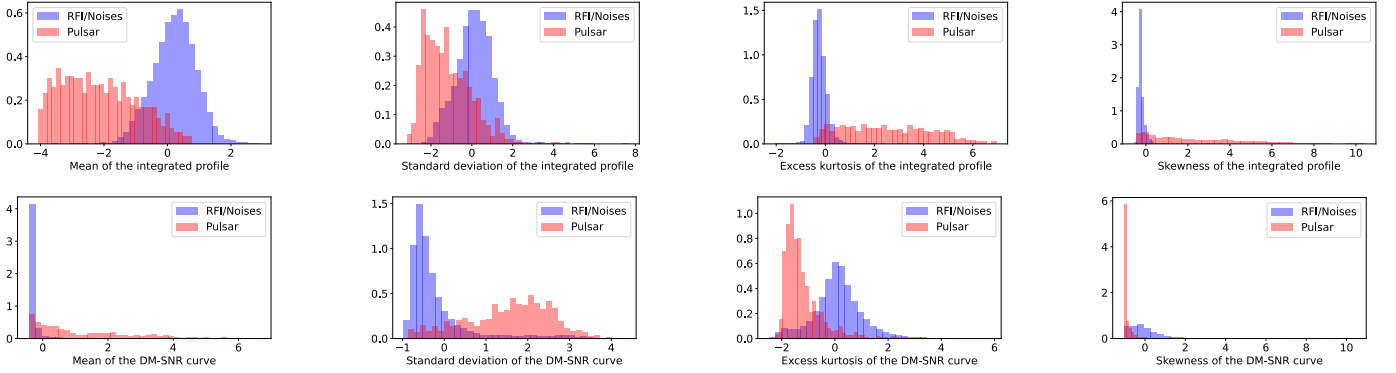


Figure 1: Z-Score Normalized Features

The feature with the *Z-Score Normalization* seem already well distributed, as shown on figure 1. Since we'll use Gaussian classifiers we can try to apply a Gaussianization over the features, but we expect worse result because of the already normalized distribution with no evident outliers.

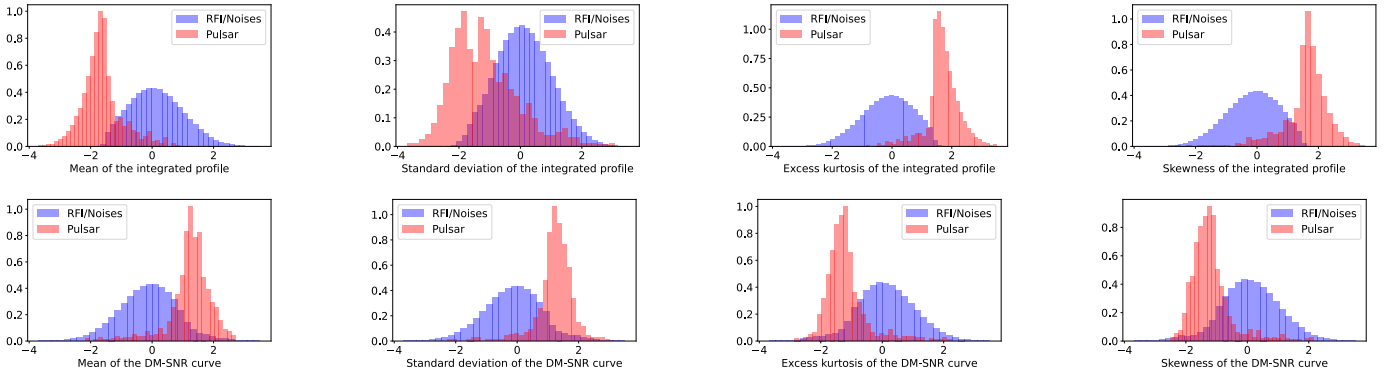


Figure 2: Gaussianized features (quantile transformation)

The heatmap over, calculated with the absolute value of the Pearson correlation coefficient on *Z-Score Normalization*

$$\left| \frac{Cov(X, Y)}{\sqrt{Var(X)}\sqrt{Var(Y)}} \right|$$

illustrates the correlation between features:




Figure 3: Heatmap - Pearson correlation

As we can see, the first 4 features (0,1,2,3) and the other 4 (4,5,6,7) are not so correlated, this can be connected to the fact that are extracted from 2 different curves (integrated profile and DM-SNR). There is strong correlation between features 6-7 and between features 2-0, a dimensionality reduction technique such as PCA could be useful to reduce features to 7 dimensions or 6, however it is possible to see that there is a moderate correlation between features 2-3 (weak for the false, strong for the true), reason why it could be possible to proceed with a PCA of 5 dimensions.

2 Classifying HTRU2 features

The approach that is followed on the following cases is *K Fold* with 3 folds, where there will be considered three applications.

- a balanced one:

$$(\tilde{\pi}, C_{fn}, C_{fp}) = 0.5, 1, 1$$

- two unbalanced:

$$(\tilde{\pi}, C_{fn}, C_{fp}) = 0.1, 1, 1$$

$$(\tilde{\pi}, C_{fn}, C_{fp}) = 0.9, 1, 1$$

Since the application target is to choose the most promising approach, we will measure performance in terms of normalized **minimum** Detection Cost Function (DCF) - it measures the cost we would pay if we made optimal decisions using the recognizer scores.

2.1 Gaussian classifiers

We start considering the MultiVariate Gaussian (MVG) classifiers: Full Covariance, Diagonal Covariance and Tied Covariance. These are generative model that assume Gaussian distributed data, given the class:

$$X|C = c \sim \mathcal{N}(\mu_c, \Sigma_c)$$

The Tied MVG classifier assumes that each class has its own μ_c , yet the covariance is shared.

$$X|C = c \sim \mathcal{N}(\mu_c, \Sigma)$$

The Diagonal, instead, assumes that the covariance matrices are diagonal.

MVG Classifiers - minDCF on the evaluation set			
3-fold			
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z-Norm features - no PCA			
MVG - Full Covariance	0.142	0.285	0.660
MVG - Diag Covariance	0.191	0.315	0.735
MVG - Tied Covariance	0.112	0.224	0.580
Z-Norm features - PCA (m = 7)			
MVG - Full Covariance	0.140	0.304	0.628
MVG - Diag Covariance	0.213	0.510	0.715
MVG - Tied Covariance	0.112	0.223	0.574
Z-Norm features - PCA (m = 6)			
MVG - Full Covariance	0.151	0.290	0.635
MVG - Diag Covariance	0.222	0.534	0.726
MVG - Tied Covariance	0.138	0.254	0.584
Gaussianized features - no PCA			
MVG - Full Covariance	0.153	0.243	0.696
MVG - Diag Covariance	0.154	0.280	0.600
MVG - Tied Covariance	0.132	0.233	0.541

We used, as said above, three different applications, and applied PCA on *Z Score Normalized* data. The Tied Covariance obtains the best performance in *K fold cross validation protocol* both with and without PCA.

It is possible to notice that PCA is not quiet effective in none of the three classifiers, but in the Tied one with m=7 we can see that it preserves the performance for our target application obtaining slightly better results for the unbalanced applications. Applying the PCA with 6 dimensions, it is also possible to notice that we have a worse performance, so we stopped at 6 dimensions.

As we have assumed before, applying Gaussianization of the features led to worse results.

The Diag model performs, in general, worse. We expected that because features are correlated in group of fours, and the Naive Bayes assumption led to bad results. In its case, the Gaussianization improve performances, but it is still the worse in all the four cases

2.2 Logistic Regression Classifiers

It is a discriminative model. We will employ two types of Logistic Regression: *Linear Logistic Regression* and *Quadratic Logistic Regression*.

Since classes are not balanced, we re-balance the costs of the different classes minimizing:

$$J(w, b) = \frac{\lambda}{2} \|w\|^2 + \frac{\pi_T}{T} \sum_{i=1|c_i=1}^n \log(1 + e^{-z_i(w^T + x_i + b)}) + \frac{1 - \pi_T}{n_F} \sum_{i=1|c_i=0}^n \log(1 + e^{-z_i(w^T + x_i + b)})$$

Where (w, b) are the model parameters. The model assumes that decision rules are linear hyper-planes orthogonal to w .

To estimate a proper value for the hyper-parameter λ we also used a 3-fold cross-validation; we consider our main application, so we set $\pi_T = 0.5$.

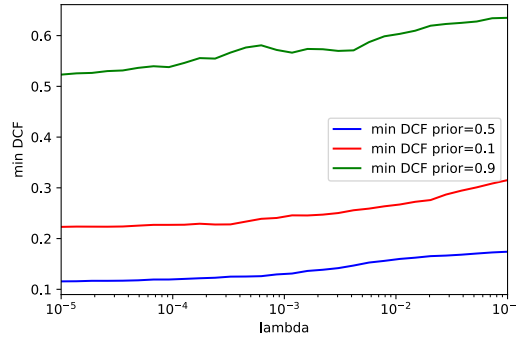


Figure 4: Linear Logistic Regression lambda tuning

As it is possible to see from the figure above, the proper value for the hyper-parameter λ would be $\lambda = 10^{-5}$. We can also consider training using a different prior π_T to see the effects on the other application, as we did over.

3-fold			
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z-Norm features - no PCA			
MVG - Tied Covariance	0.112	0.224	0.580
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.116	0.223	0.523
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.1$)	0.114	0.217	0.547
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.9$)	0.123	0.226	0.523
Z-Norm features - PCA (m = 7)			
MVG - Tied Covariance	0.112	0.223	0.574
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.116	0.220	0.532
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.1$)	0.115	0.218	0.551
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.9$)	0.120	0.222	0.541
Gaussianized features - no PCA			
MVG - Tied Covariance	0.132	0.233	0.541
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.124	0.232	0.517
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.1$)	0.128	0.230	0.519
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.9$)	0.127	0.235	0.510

We tested the application first with *Z-Score Normalization*, without applying PCA, data, then with *Z-Score Normalization* with a PCA of 7 dimensions applied, and at the end with *Gaussianized features*.

Overall, the MVG Tied Covariance performs better than the logistic regression.

Balancing the training priors, we can see that the Logistic Regression performs better than the MVG Tied.

surprisingly Gaussianized features does improve the unbalance application $\pi = 0.9$, even if the *Logistic Regression* does not require specific assumption on data distribution.

We will now focus on the Quadratic Logistic Regression. We still have to tune a value for the hyper-parameter λ :

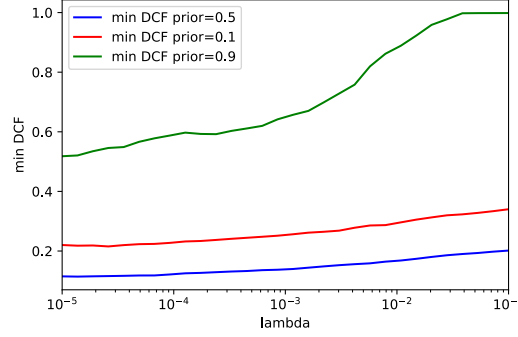


Figure 5: Linear Logistic Regression lambda tuning

Also in this case, the selected value will be $\lambda = 10^{-5}$.

3-fold			
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z-Norm features - no PCA			
MVG - Tied Covariance	0.112	0.224	0.580
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.116	0.223	0.523
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.1$)	0.114	0.217	0.547
Quad Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.115	0.220	0.518
Quad Log Reg($\lambda=10^{-5}$, $\pi_T = 0.1$)	0.114	0.213	0.553
Quad Log Reg($\lambda=10^{-5}$, $\pi_T = 0.9$)	0.118	0.232	0.558
Z-Norm features - PCA (m = 7)			
MVG - Tied Covariance	0.112	0.223	0.574
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.116	0.220	0.532
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.1$)	0.115	0.218	0.551
Quad Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.113	0.220	0.532
Quad Log Reg($\lambda=10^{-5}$, $\pi_T = 0.1$)	0.113	0.215	0.546
Quad Log Reg($\lambda=10^{-5}$, $\pi_T = 0.9$)	0.117	0.231	0.548
Gaussianized features - no PCA			
MVG - Tied Covariance	0.132	0.233	0.541
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.124	0.232	0.517
Quad Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.128	0.232	0.520

As we did for the *Linear Logistic Regression*, we try with different π_T .

We can notice that the *Quadratic Logistic Regression* performs worse than the *Linear Logistic Regression*.

Things change, tho, if we apply PCA with 7 dimensions. In fact, also if the MVG Tied Covariance is still the best, we can see that the *Quadratic Logistic Regression* performs similar to it, gaining better results in the unbalanced applications.

Using different values of π_T improves the classification performance for the specific application (at least for $\pi = 0.1$)

Gaussianizing features is not relevant, as we expected.

2.3 Support Vector Machines

We now turn our attention to SVMs, where we will consider:

- Linear Support Vector Machine
- Polynomial quadratic kernel
- Radial Basis Function kernel formulations

The formulation of SVM can be expressed as:

$$\max_{\alpha} \alpha^T \mathbf{1} - \frac{1}{2} \alpha^T \hat{H} \alpha$$

Subject to:

$$0 \leq \alpha_i \leq C_i, i = 1, \dots, n$$

where n is the number of training samples, \mathbf{C} is a hyper-parameter that we need to tune, $\mathbf{1}$ is a n -dimensional vector of ones, z_i is the class label for the i -th sample encoded as:

$$z_i = \begin{cases} +1, & \text{if } x_i \text{ belongs to class 1 (true pulsar)} \\ -1, & \text{if } x_i \text{ belongs to class 0 (false pulsar)} \end{cases} \quad (1)$$

and \mathbf{H} is a matrix whose elements are

$$H_{i,j} = z_i z_j x_i^T x_j$$

Since we have an hyper-parameter, we need to tune and choose the best value. We used again a 3-fold cross-validation with:

- Z-Score Normalization
- Z-Score Normalization and PCA with 7 dimensions
- Gaussianized features

We will first analyze the *Linear Support Vector Machine*.




Figure 6: Linear SVM C tuning: Z - Score, Z - Score with PCA 7, Gaussianized features

Given the analysis of the figure above, we choose $C = 0.1$ since it seems to provide a lower minDCF along our main application (blue line). Chosen C , we can analyze the results that we obtained with a 3-fold cross-validation.

	3-fold		
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z-Norm features - no PCA			
MVG - Tied Covariance	0.112	0.224	0.580
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.116	0.223	0.523
Quad Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.115	0.220	0.518
Linear SVM ($C = 0.1$)	0.114	0.218	0.563
Z-Norm features - PCA (m = 7)			
MVG - Tied Covariance	0.112	0.223	0.574
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.116	0.220	0.532
Quad Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.113	0.220	0.532
Linear SVM ($C = 0.1$)	0.113	0.218	0.556
Gaussianized features - no PCA			
MVG - Tied Covariance	0.132	0.233	0.541
Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.124	0.232	0.517
Quad Log Reg($\lambda=10^{-5}$, $\pi_T = 0.5$)	0.128	0.232	0.520
Linear SVM ($C = 0.1$)	0.134	0.240	0.527

Linear SVM produces slightly worse results than MVG. With PCA m=7 it aligns with MVG and Quadratic LogReg.

As we expected, Gaussianization does not improve our model.

We can now focus on the **Non-Linear Support Vector Machines**. The formulation of SVM can be also expressed as:

$$\hat{x} = \begin{bmatrix} X \\ K \end{bmatrix}$$

It is needed then to find a function $k(\hat{x}_i, \hat{x}_j) = \phi(\hat{x}_i)^T \phi(\hat{x}_j)$ that allows us to compute the dot product. We can thus define:

$$k(\hat{x}_i, \hat{x}_j) = (\hat{x}_i^T \hat{x}_j + c)^d$$

We will employ $d = 2$ for *Quadratic Support Vector Machine* and tune again to find an optimal value for the hyper-parameter C .

The technique used to estimate C is still a 3-fold cross-validation.




Figure 7: Quadratic SVM C tuning: Z - Score, Z - Score with PCA 7, Gaussianized features

Given the results of the figure above, the choice of C isn't too critical, we can employ $C = 0.1$.

Chosen C , we can again see the results of the calculation:

	3-fold		
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z-Norm features - no PCA			
Linear SVM ($C = 0.1$)	0.114	0.218	0.563
Quadratic SVM($C=0.1$)	0.117	0.212	0.586
Z-Norm features - PCA ($m = 7$)			
Linear SVM ($C = 0.1$)	0.113	0.218	0.556
Quadratic SVM($C=0.1$)	0.117	0.215	0.592
Gaussianized features - no PCA			
Linear SVM ($C = 0.1$)	0.134	0.240	0.527
Quadratic SVM($C=0.1$)	0.135	0.220	0.603

We can notice that, confronted to the *Linear SVM*, the **Quadratic SVM** performs slightly worse, in fact its minDCF is 0.117 in both application with and without PCA with 7 dimensions. As seen before, the *Linear SVM* slightly improve with the applicaiton of the PCA.

When we try to Gaussianize features, instead, we obtain a worse result, even if *Linear SVM* and *Quadratic SVM* are almost the same: in fact the *Linear SVM* obtains a minDCF of 0.134, on the other hand the *Quadratic SVM* obtains a minDCF of 0.135.

We can also use a *different* kernel function:

$$k(\hat{x}_i, \hat{x}_j) = e^{-\gamma \|\hat{x}_i - \hat{x}_j\|}$$

That is also called **Radial Basis Function Support Vector Machine**.

We tune again for the best value of the hyper-parameters C and γ .




Figure 8: RBF SVM C tuning

As Figure 8 shows, good values for the model could be $C = 10$ and $\gamma = 0.01$. For Gaussianize features we choose $\gamma = 0.1$ instead.

At this point we can again see the results for the **Radial Basis Function**.

3-fold			
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z-Norm features - no PCA			
Linear SVM ($C = 0.1$)	0.114	0.218	0.563
Quadratic SVM($C=0.1$)	0.117	0.212	0.586
RBF SVM($C=10, \gamma=0.01$)	0.113	0.210	0.553
Z-Norm features - PCA ($m = 7$)			
Linear SVM ($C = 0.1$)	0.113	0.218	0.556
Quadratic SVM($C=0.1$)	0.117	0.215	0.592
RBF SVM($C=10, \gamma=0.01$)	0.112	0.209	0.564
Gaussianized features - no PCA			
Linear SVM ($C = 0.1$)	0.134	0.240	0.527
Quadratic SVM($C=0.1$)	0.135	0.220	0.603
RBF SVM($C=10, \gamma=0.1$)	0.127	0.215	0.798

It is possible to see that the **RBF**, compared to the *Linear* and *Quadratic Support Vector Machine* slightly improve in both the application, with and without applying the PCA with a reduction to 7 dimensions.

Gaussianizing features, we can notice that there are worse performances but the **RBF** would still be the better to perform.

The **Radial Basis Function Support Vector Machine** seems to be the best model among the *SVMs*, anyway we can try to rebalance the classes, using the dual formulation subject to:

$$0 \leq \alpha_i \leq C_i, \quad i = 1, \dots, n$$

Where $C_i = C_T$ for samples of class 1 and $C_i = C_F$ for samples of class 0. We select $C_T = C \frac{\pi_T}{\pi_T^{emp}}$ and $C_F = C \frac{\pi_F}{\pi_F^{emp}}$.

For $\pi_T = 0.5$ we obtain:

	3-fold		
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z-Norm features - PCA (m = 7)			
Linear SVM (C = 0.1)	0.113	0.218	0.556
Quadratic SVM(C=0.1)	0.117	0.215	0.592
RBF SVM(C=10, $\gamma=0.01$)	0.112	0.209	0.564
RBF SVM(C=10, $\gamma=0.01$, $\pi_T=0.5$)	0.110	0.220	0.528

RBF with $\pi_T = 0.5$ obtains slightly better results than the previous best models (MVG and Quadratic LogReg).

2.4 Gaussian Mixture Models

The last model considered is a Gaussian Mixture Model (GMM) over the data of each class. The hyperparameter to tune, in this case, is the number of Gaussian components. It will be tuned both for *Z-Score Normalized* and *Gaussianized* features.

In this report we will consider:

- Diagonal covariance model
- Full covariance model
- Tied covariance model
- Tied Diagonal covariance model

We will use again the 3-folds approach to choose the number of components and to compare different models.




Figure 9: Number of components for GMM models. x is the number of slits. $\pi = 0.5$

The GMM model that performs better is the Full Covariance, in fact is possible to see that with a number of 3 splits, so $2^3 = 8$ components, the minDCF is the lowest, reaching a value of 0.113.

	3-fold		
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z-Norm features			
GMM - Full Covariance (8 components, Z Score Normalization)	0.113	0.235	0.559
GMM - Tied Covariance (32 components, Z Score Normalization)	0.135	0.244	0.640
GMM - Diagonal Covariance (32 components, Z Score Normalization)	0.147	0.266	0.645
GMM - Tied Diagonal Covariance (64 components, Z Score Normalization)	0.149	0.298	0.628

In the case of the GMM model the Tied is different because the Tied MVG assumes that the covariance matrix is shared between class; yet in the case of a GMM the covariance matrix is shared between *components* of the same class.

Based on Figure 9, we selected the best four models:

- GMM - Full Covariance with 8 components and Z Score Normalization
- GMM - Tied Covariance with 32 components and Z Score Normalization
- GMM - Diagonal Covariance with 32 components and Z Score Normalization
- GMM - Tied Diagonal Covariance with 64 components and Z Score Normalization

Analyzing these results, it is possible to notice that the best model appears to be the Full Covariance GMM with 8 components and *Z-Score Normalized* features, without applying PCA.

Now we consider the three model that obtains similar results:

- MVG Tied PCA (m=7)
- Quadratic LogReg(= 10^{-5} , $\pi_T = 0.5$) PCA (m=7)
- RBF SVM($C = 10$, $\gamma = 0.01$, $\pi_T = 0.5$) PCA (m=7)

for this application we will select two candidate: RBF SVM($C = 10$, $\gamma = 0.01$, $\pi_T = 0.5$) PCA (m=7) as a first system, and MVG Tied PCA (m=7) as a secondary system. Quadratic LogReg could be a substitute for the MVG because we have seen it gives better results on unbalanced application but we will not consider it.

2.5 Score Calibration

We have discussed about minDCF metric and costs for the *validation set* using scores of the recognizer. However, the cost we actually pay, depends on a threshold we use to perform class assignment, reason why we introduce the **actualDCF**.

We decided to follow an approach based on re-calibrating the scores, transforming them so that the theoretical threshold

$$t = -\log \frac{\tilde{\pi}}{1 - \tilde{\pi}}$$

provides close to optimal values over a wide range of effective priors $\tilde{\pi}$.

In this approach we need a transformation function f that maps the classifiers scores into well-calibrated scores $s_{cal} = f(s)$.

There are more than one way to estimate f , such as:

- Isotonic regression
- Discriminative score models (such as Logistic Regression)
- Generative score models

We will focus on the second approach, in fact our function f will have a simple form such as:

$$f(s) = \alpha s + \beta$$

$f(s)$ can be thus interpreted as the log-likelihood ratio for two class hypotheses:

$$f(s) = \log \frac{f_{S|C}(s|H_T)}{f_{S|C}(s|H_F)} = \alpha s + \beta$$

The class posterior for $\tilde{\pi}$ corresponds then to:

$$\log \frac{P(C = H_T|s)}{P(C = H_F|s)} = \alpha s + \beta + \log \frac{\tilde{\pi}}{1 - \tilde{\pi}}$$

Interpreting scores as features, we have a quiet similar expression as the log posterior ratio of the logistic regression model.

In fact, if we let

$$\beta' = \beta + \log \frac{\tilde{\pi}}{1 - \tilde{\pi}}$$

we have the exact same model. This, can be employed as a model parameter over our training scores. To recover the calibrated score $f(s)$ we will need then:

$$f(s) = \alpha s + \beta = \alpha s + \beta' - \log \frac{\tilde{\pi}}{1 - \tilde{\pi}}$$

For the **SVM RBF** with calibrated scores, the results are:

	minDCF ($\pi = 0.5$)	minDCF $\pi = 0.1$)	minDCF ($\pi = 0.9$)
	0.110	0.220	0.528
	actDCF ($\pi = 0.5$)	actDCF ($\pi = 0.1$)	actDCF ($\pi = 0.9$)
Log Reg ($\pi_T = 0.5$)	0.112	0.225	0.558

Meanwhile for the **MVG Tied**, the results with calibrated scores are:

	minDCF ($\pi = 0.5$)	minDCF ($\pi = 0.1$)	minDCF ($\pi = 0.9$)
	0.112	0.223	0.574
	actDCF ($\pi = 0.5$)	actDCF ($\pi = 0.1$)	actDCF ($\pi = 0.9$)
Log Reg ($\pi_T = 0.5$)	0.113	0.231	0.604

We selected these two algorithms because their results showed the best performances. Both *minDCF* are with PCA with 7 dimensions applied.

We can see from the tables above that the **actualDCF** is slightly higher than the minDCF, due to the calibration of the scores.




Figure 10: Calibrated (left) and non calibrated scores (right) of SVM and MVG

We can observe that the Logistic Regression approach obtains accurate results. It is also possible to see from the Bayes error plots for the calibrated and non-calibrated scores, the model is trained with $\tilde{\pi} = 0.5$.

2.5.1 Fusion

Two classifiers can agree and disagree on some decision, we can try to combine the decisions at **score-level**, fusing the output of the two classifiers to train a meta-model over the scores.

The technique that we use is the stacking ensemble: we take the output of the two classifiers (scores), we build a new dataset where each feature is the output of each classifier and then we train a meta-model (in our case a balanced LogReg) to produce new score balanced with the model parameters.

We perform a fusion of the **SVM** and **MVG**, hopefully the new scores will have higher accuracy and benefit from the re-calibrating effect of the model.




Figure 11: ROC and DET plots

	$\pi = 0.5$		$\pi = 0.1$		$\pi = 0.9$	
	minDCF	actDCF	minDCF	actDCF	minDCF	actDCF
SVM	0.110	0.112	0.220	0.225	0.528	0.558
MVG	0.112	0.113	0.223	0.231	0.574	0.604
Fusion	0.109	0.109	0.218	0.220	0.540	0.560

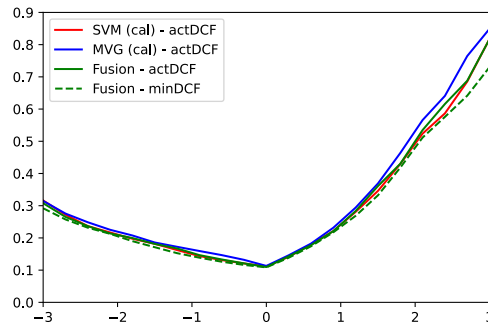


Figure 12: Bayes Error plot, fused system

Fusion does not really improve the scores, it still reduce very slightly the costs but it's aligned with the original models.

Scores are well-calibrated.

We expect the fusion to obtain results close to the RBF SVM or lower.

3 Experimental results

We now focus on analyzing the different models performances on the test set, in terms of minDCF_s, to see how our choices affected performance for unseen data.

MVG models:

All data			
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z - Norm features			
MVG - Full Cov	0.140	0.283	0.646
MVG - Tied Cov	0.110	0.207	0.591
MVG - Diag Cov	0.185	0.330	0.621
Z-Norm features - PCA (m = 7)			
MVG - Full Cov	0.139	0.293	0.574
MVG - Tied Cov	0.110	0.208	0.587
MVG - Diag Cov	0.201	0.514	0.755
Z-Norm features - PCA (m = 6)			
MVG - Full Cov	0.139	0.293	0.574
MVG - Tied Cov	0.110	0.208	0.587
MVG - Diag Cov	0.201	0.514	0.755
Gaussianized features			
MVG - Full Cov	0.151	0.241	0.622
MVG - Tied Cov	0.125	0.221	0.540
MVG - Diag Cov	0.154	0.285	0.574

LR models:

All data			
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z - Norm features			
LR ($\lambda=1e-5$, $\pi_T=0.5$)	0.109	0.199	0.537
LR ($\lambda=1e-5$, $\pi_T=0.1$)	0.109	0.197	0.529
LR ($\lambda=1e-5$, $\pi_T=0.9$)	0.115	0.205	0.508
Z-Norm features - PCA (m = 7)			
LR ($\lambda=1e-5$, $\pi_T=0.5$)	0.108	0.203	0.538
LR ($\lambda=1e-5$, $\pi_T=0.1$)	0.109	0.199	0.532
LR ($\lambda=1e-5$, $\pi_T=0.9$)	0.115	0.203	0.532
Gaussianized features			
LR ($\lambda=1e-5$, $\pi_T=0.5$)	0.121	0.216	0.485
LR ($\lambda=1e-5$, $\pi_T=0.1$)	0.122	0.209	0.499
LR ($\lambda=1e-5$, $\pi_T=0.9$)	0.121	0.217	0.475

Quadratic LR:

All data			
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z - Norm features			
Quad LR ($\lambda=1e-5$, $\pi_T = 0.5$)	0.103	0.203	0.470
Quad LR ($\lambda=1e-5$, $\pi_T = 0.1$)	0.105	0.202	0.519
Quad LR ($\lambda=1e-5$, $\pi_T = 0.9$)	0.108	0.207	0.451
Z-Norm features - PCA (m = 7)			
Quad LR ($\lambda=1e-5$, $\pi_T = 0.5$)	0.104	0.202	0.470
Quad LR ($\lambda=1e-5$, $\pi_T = 0.1$)	0.104	0.202	0.515
Quad LR ($\lambda=1e-5$, $\pi_T = 0.9$)	0.108	0.208	0.467
Gaussianized features			
Quad LR ($\lambda=1e-5$, $\pi_T = 0.5$)	0.116	0.215	0.479
Quad LR ($\lambda=1e-5$, $\pi_T = 0.1$)	0.119	0.209	0.507
Quad LR ($\lambda=1e-5$, $\pi_T = 0.9$)	0.116	0.223	0.483

Support Vector Machines:

All data			
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
Z - Norm features			
SVM C=0.1	0.112	0.205	0.552
QSVM C=0.1	0.102	0.204	0.522
RBF C=10, $\gamma=0.01$	0.106	0.200	0.465
Z-Norm features - PCA (m = 7)			
SVM C=0.1	0.112	0.205	0.552
QSVM C=0.1	0.104	0.204	0.532
RBF C=10, $\gamma=0.01$	0.106	0.201	0.483
RBF C=10, $\gamma=0.01$, $\pi_T = 0.5$	0.098	0.200	0.480
Gaussianized features			
SVM C=0.1	0.129	0.222	0.499
QSVM C=0.1	0.125	0.209	0.541
RBF C=10, $\gamma=0.1$	0.114	0.209	0.687

GMM models:

	All data		
	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
	Z - Norm features		
GMM Full-Cov 8c	0.100	0.220	0.578
GMM Tied 32c	0.128	0.241	0.607
GMM Diag 32c	0.142	0.267	0.616
GMM Tied-Diag 64c	0.145	0.282	0.683

We can see that the results are consistent with the one showed during the *validation*, with a few exception: in fact the LogReg models here have better performance than the MVG Tied model (candidate system that we selected).

GMM Full-Cov 8c it's the second best model for this experiment.

Overall PCA (m=7) in some cases is effective but in others it worsen the results (Quad SVM).

The first system that we selected is still the best one.

Finally, we can evaluate the effectiveness of our model fusion.

Fusion:

	$\pi = 0.5$		$\pi = 0.1$		$\pi = 0.9$	
	minDCF	actDCF	minDCF	actDCF	minDCF	actDCF
SVM - Calibrated	0.098	0.099	0.200	0.210	0.480	0.499
MVG - Calibrated	0.110	0.114	0.208	0.220	0.587	0.606
Fusion	0.098	0.099	0.199	0.206	0.494	0.534




Figure 13: ROC and DET plots on evaluation set

As we expected, the fusion aligns with the SVM model, giving similar results.

4 Conclusion

Concluding, our approach, consisting of the fusion of two subsystems, it's effective: we could've achieved better results fusing different system, but the one we analyzed produce the best results among all the models we made and produce well calibrated scores.

We are able to achieve a DCF cost of ≈ 0.1 for our main application ($\pi = 0.5$). The model is good also for applications with imbalanced prior 0.1 (DCF cost of ≈ 0.2) and for the one with prior 0.9. (DCF cost of ≈ 0.5).

The choices we made on our training / validation sets proved effective also for the evaluation data

References

- [1] R. J. Lyon, B. W. Stappers, S. Cooper, J. M. Brooke, J. D. Knowles, Fifty Years of Pulsar Candidate Selection: From simple filters to a new principled real-time classification approach, Monthly Notices of the Royal Astronomical Society 459 (1), 1104-1123, DOI: 10.1093/mnras/stw656
- [2] R. J. Lyon, HTRU2, DOI: 10.6084/m9.figshare.3080389.v1.